



Storage virtualization and the HP StorageWorks Enterprise Virtual Array

Table of contents

Abstract.....	2
Introduction.....	2
What is storage virtualization?.....	2
Server level	3
Storage network or SAN fabric level	3
Storage system level	4
HP StorageWorks virtualization implementation at the storage system level	4
Improved performance.....	4
Simplified management	5
Dynamic expansion.....	6
Distributed sparing	7
Data replication tools	8
Traditional snapshot	8
Virtually Capacity-Free Snapshot	8
Virtually Instantaneous Snapclone.....	8
EVA remote replication with HP StorageWorks Continuous Access EVA.....	9
Virtualization and tiered storage.....	9
The business value of virtualization	10
Time	10
Cost	10
Space.....	11
Conclusion.....	11
For more information.....	13

Abstract

This white paper presents an overview of storage virtualization at the controller level of a storage system. Virtualization technology allows companies to respond quickly to ever-changing storage capacity requirements. Virtualization has multiple definitions and many different implementations in a storage area network (SAN). This document describes how HP defines virtualization technology and how it applies to the HP StorageWorks Enterprise Virtual Array (EVA) storage system. It also explores how virtualization of storage provides advantages over conventional storage through performance, management, distributed sparing, Virtually Capacity-Free Snapshot replication, and Virtually Instantaneous Snapclone replication.

This document explains tiered storage and the use of low-cost Fibre Attached Technology Adapted (FATA) drives. It illustrates the cost benefit of FATA drives and introduces the concept of associating the cost of the storage to the value of the data stored on it. Finally, it discusses how the EVA can save a company time, space, and total cost of ownership through the use of virtualization and the advantages it provides to firms doing business in today's climate.

Introduction

IT professionals are always looking for better ways to access and protect data. Organizations are seeking technologies that reduce costs and enhance the management of their rapidly growing storage requirements. Businesses have implemented technologies such as RAID, tape back-ups, and SANs in an attempt to control large amounts of data for business-critical operations. As today's companies face even greater challenges accessing and managing their business-critical information, the need for more effective solutions has become a business necessity.

Storage virtualization is one of the most exciting and beneficial solutions for storage management. This paper presents an overview of the different layers of virtualization, why this technology has tremendous merit in the area of networked storage, the HP implementations of virtualization, and the significant benefits provided by this approach to data storage.


What is storage virtualization?

Storage virtualization is generally defined as the "transparent abstraction of storage at the block level." In essence, virtualization separates logical data access from physical data access, enabling users to create large storage pools from physical storage. Virtual disks are created from these pools and are allocated to servers on the network as logical storage when needed.

Virtual storage reduces the physical one-to-one relationship between servers and storage devices. The physical disk devices and distribution of storage capacity become transparent to servers and applications. According to John Webster, a storage industry analyst with Illuminata, "The single most important attribute of any storage virtualization solution is the ability to mask complexity and thereby make manageable that which is increasingly becoming unmanageable." The goal of products and solutions using storage virtualization therefore is to simplify management.

There are three levels within a networked storage environment in which virtualization can occur: the server level, storage network or SAN fabric level, and storage system level. These levels can be used together or independently to increase the benefits to users substantially. The following sections define these virtualization levels in more detail.

Figure 1. Implementing Storage Virtualization



Server level

At the server level, virtualization can be implemented through software residing on the server and is independent of storage devices. With this software, the operating system causes the server to behave as if it is in communication with a device type even though it is actually communicating with a virtual disk. Server-based virtualization can be deployed in homogeneous SAN and non-SAN environments. It has limited interoperability with hardware or software components.

As an example of server-based storage virtualization, HP offers HP PolyServe Software for Microsoft SQL Server. This product enables multiple Microsoft SQL Server instances to be consolidated onto substantially fewer servers and centralized SAN storage. HP PolyServe's Dynamic Re-Hosting™ allows you to deploy, stack, or reallocate Microsoft SQL Server instances to any server quickly. It also enables you to add server and storage capacity easily and without disrupting existing operations.

Storage network or SAN fabric level

The HP vision supports the creation of an open SAN—an infrastructure based on industry standards that delivers the "many-to-many" functionality needed to meet the storage requirements of today's computing environments. This functionality includes scaling, virtualization, automation, simplification, interoperability, and investment protection.

SANs provide the data communication infrastructure needed for the most advanced and most cost efficient computer mass storage systems. SANs provide unprecedented levels of flexibility in system management and configuration. Servers can be added to and removed from a SAN while their data remains in the SAN. Multiple servers can access the same storage for more consistent and rapid processing. The storage itself can be easily increased, changed, or re-assigned.

SANs reduce total cost of ownership for both large and small storage systems. SAN technology supports the management features and I/O price performance demanded in today's competitive IT environment and offers the storage component investment protection needed to reduce capital expense.

Storage system level

At the storage system level, virtualization is implemented on storage array controllers, independent of the host. The storage array controllers can create virtual disks, snapshots, and clones in conjunction with management software. This storage virtualization process is managed centrally using a storage management server and a Web browser.

Virtualization in the individual storage system controller is the evolutionary next step beyond classic RAID technology. Storage virtualization at this level is ideal for environments requiring high performance, high data availability, fault tolerance, efficient storage management, data replication, and cluster support.

The HP StorageWorks Enterprise Virtual Array (EVA) is a premier example of the advantages of virtualization at the storage system level. The EVA provides significant benefits in automation and ease of use for customers. The following sections describe in detail how the EVA provides these benefits.

Table 1. Virtualization at various levels of the SAN

Server level	Ideal for small and entry-level storage networking
Storage network or SAN fabric level	SAN-wide virtualization techniques that increase efficiency in SAN development, management, and service are beneficial to large heterogeneous SAN environments
Storage system level	Storage virtualization for large capacity environments that provides benefits such as Virtually Capacity-Free Snapshot for any Vraid type, allowing allocation of capacity on demand, and does not sacrifice performance or reliability

HP StorageWorks virtualization implementation at the storage system level

HP defines virtualization at the storage system level as the disassociation of capacity from its underlying physical disk restrictions to create large pools of storage available to multivendor hosts. These pools of virtual capacity can be configured as virtual disks and presented to any or all connected hosts. This implementation of virtualization allows for more efficient use of storage capacity, simplified management, and an overall reduction in costs.

The EVA provides tremendous benefits as a result of virtualization. Along with meeting the standard requirements such as high availability, data protection, and fault tolerance, the EVA also provides many value-added features to help organizations increase their storage environments. This technology capitalizes on virtualization to provide optimal performance, ease of management, improved capacity utilization, powerful data replication tools, and faster rebuild times, while simplifying the management of the virtualized storage capacity. The following sections focus on the value-added benefits provided by the EVA and how this system truly capitalizes on virtualization technology.

Improved performance

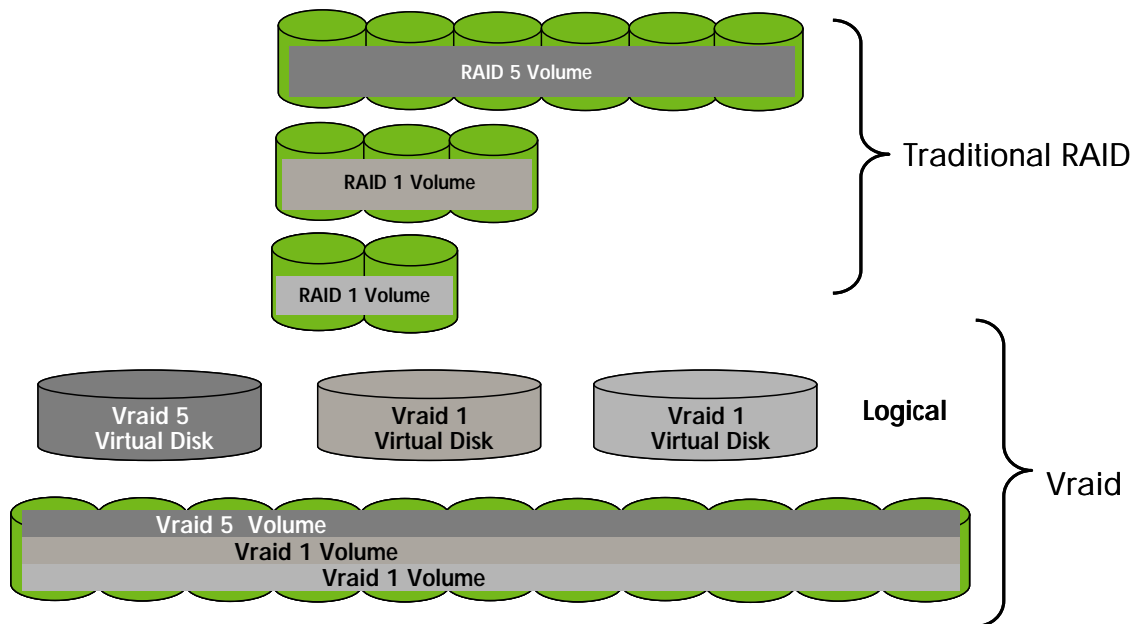
HP believes that performance is one of the most important benefits derived from virtualization at the storage system level. The main reason for this performance improvement is load balancing through the elimination of hot spindles, or hot drives. A “hot spindle” is a drive that is subjected to intense I/O loads because the data it contains is frequently accessed or it is heavily worked as a result of the type of RAID employed. Virtualization allows data to be spread across more spindles (disks), dramatically improving performance. Disks no longer need to be arranged in traditional RAID sets by disk capacity size, disk type (high performance or high capacity), and type of RAID protection. Rather, a “virtual”

disk draws its capacity from a designated pool of storage. Because of virtualization-enabled storage pooling, the EVA can support multiple virtual disks of varying capacity and RAID types within a single storage pool. In addition, all virtual disks in a pool will spread their capacity across all the physical disks that contribute to that pool.

For example, a traditional RAID 1 (mirrored) configuration for 146 GB of data space requires two 146 GB disks to be assigned as RAID 1 (see Figure 2). When the data, which might be in a high performance application, is read or written, only two disks are involved, thus limiting performance. With virtualization, this same 146 GB RAID 1 requirement is spread across four, six, eight, or more disks in a pool. Performance can be greatly improved because more spindles are simultaneously involved in larger transfers.

Virtualization allows data to be redistributed across physical disks within a storage pool if an activity occurs that causes a change to the virtual disk data or pool structure. As a result, the EVA utilizes an on-the-fly leveling algorithm to balance performance without interrupting ongoing workloads. This process redistributes the blocks of each virtual disk evenly across as many spindles as the redundancy type (Vraid level) of the virtual disk will allow. This leveling process is activated whenever the EVA detects an opportunity to improve utilization, such as a change in the number of disks in the pool.

Figure 2. Traditional RAID comparison with Vraid with EVA



Simplified management

Simplified storage management is another key benefit of virtualization at the storage system level. It enables administrators to handle the attributes of storage rather than the mechanics of storage management. Virtualization enables users to create a single management model for their storage regardless of RAID type, thus eliminating manual data placement tasks.

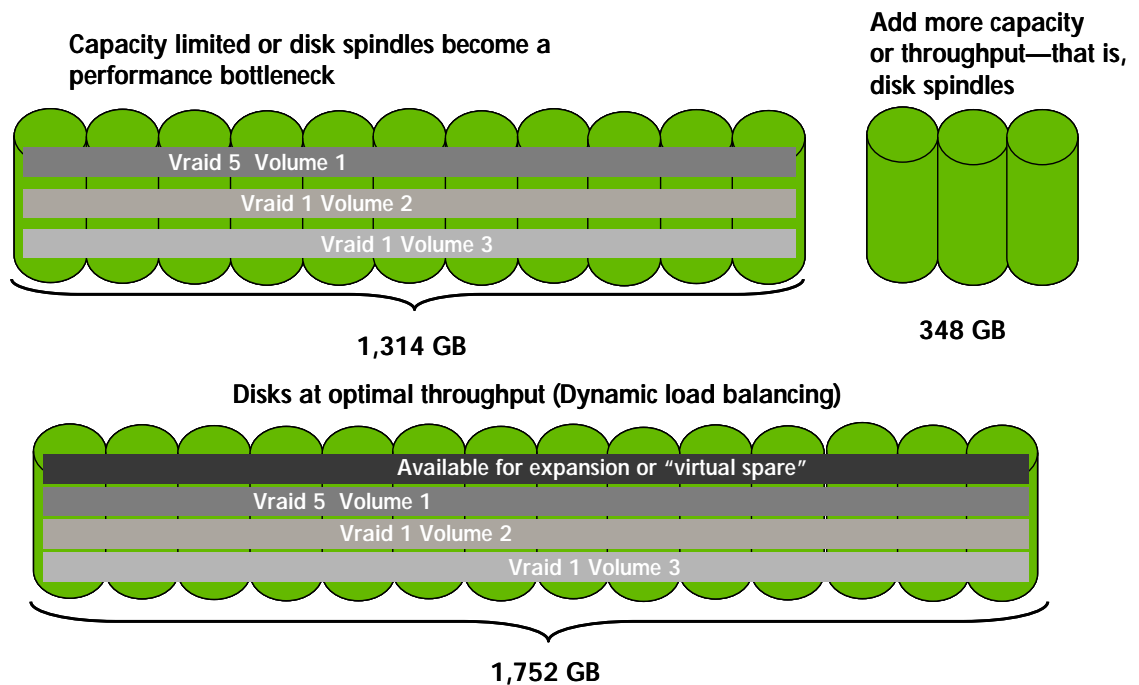
This reduces the complexity of storage deployment significantly, enables administrators to manage storage assets as a consolidated pool, and shifts capacity management from the individual unit to the pool level. All these functions simplify storage management, improve capacity utilization, and reduce training requirements.

Through the use of HP StorageWorks Command View EVA on an application server or a Storage Management Appliance, administrators can manage up to 16 EVAs on the SAN. This centralized management point and Web browser capability from other servers enables an administrator to manage significantly more storage capacity than in previous architectures. The ease of use, high disk utilization, and compact footprint of the EVA provides an ideal solution for IT consolidation.

Dynamic expansion

The ability to expand or shrink virtual disk capacity dynamically—without application downtime—improves efficiency at the system level greatly. Virtualization enables administrators using the EVA to monitor the capacity usage of a volume or the storage pool and dynamically allocate or remove capacity to either as needed. This functionality eliminates “stranded capacity” seen in other architectures. Stranded capacity is capacity that is pre-allocated to a volume to account for the growth of the volume over time. The ease of volume expansion with the EVA eliminates this traditional requirement, providing significantly better capacity utilization over traditional architectures.

Figure 3. Capacity or throughput expansion with EVA



For example, consider a 500 GB virtual RAID 5 (Vraid 5) set. As user files increase in size, capacity approaches the 500 GB limit of the volume. At this point, an administrator can add capacity to the existing pool, and the new capacity will be brought online dynamically (see Figure 3). This process can be performed automatically or manually as directed by the storage administrator. When capacity allocation occurs, the virtual disk will expand and the existing data will be releveled across all of the spindles in the volume as a background task—all with zero downtime and no performance impact.

This ability for the user to increase the virtual disk size (subject to available free capacity in the pool) results in outstanding gains in capacity usage and consolidation. This feature of virtualization allows storage use to be more efficient by eliminating stranded capacity and redistributing data as the pool grows. The result is that you purchase less initial storage capacity and less over the long-term.

Distributed sparing

Distributed or virtual sparing of disk capacity provides faster rebuild times in the EVA, which is a benefit because it reduces the possible vulnerability of data to an additional failed drive. In traditional architectures, storage administrators must set aside a physical drive as a dedicated “hot” spare in case of a drive failure. Because of its virtualizing architecture, the EVA distributes spare capacity in the same manner as a dedicated spare, across multiple spindles. This functionality allows data to be stored on each spindle, further balancing the I/O load across additional disks.

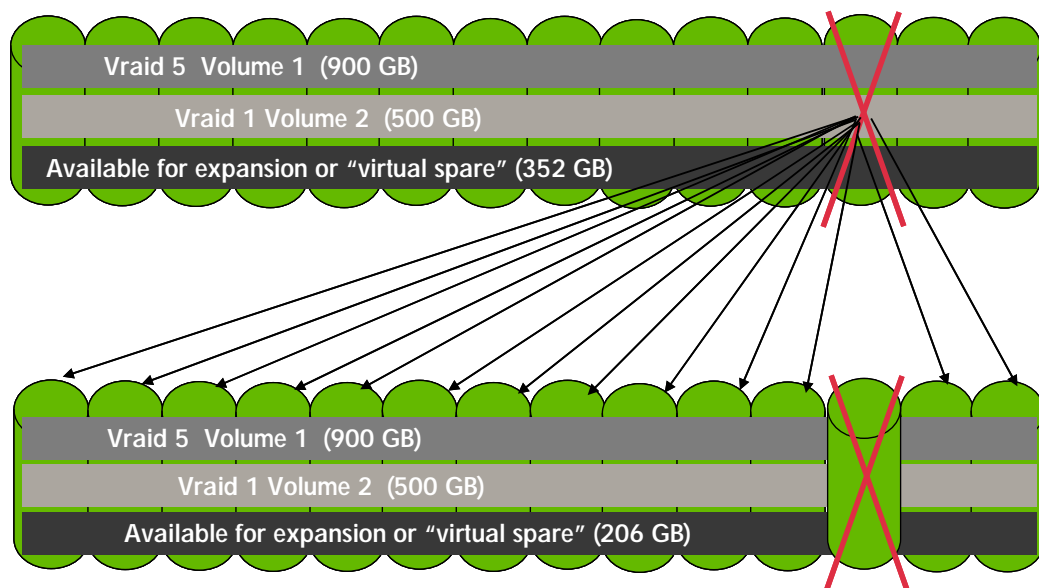
The most obvious benefit for carrying spare capacity in this manner is faster rebuild times in the case of a drive failure. Because the data is striped across multiple spindles, the EVA can rebuild data faster because more spindles are accessing data (see Figure 4).

Management software can be used to hold back enough pool space to spare an optimal number of drives per pool, depending on the application. This spare space is used to restore redundancy to virtual disks that have segments mapped to a disk drive that fails. If a drive fails, the EVA initiates a background process that rebuilds Vraid 1 or Vraid 5 data into the unused space of the pool.

The EVA first utilizes any unused pool space for sparing and if there is no unused capacity available (either because of normal usage or previous drive failures) then the distributed spare space would be used. In this manner the EVA enhances availability by using all available capacity, unused pool space and designated spare space, to mitigate disk failures.

Distributed sparing works with an algorithm similar to load leveling, except the goal is to empty the failed or failing drive as quickly as possible, rather than level its usage with other drives.

Figure 4. Distributed Sparing



Data replication tools

The EVA controller provides three types of local data replication tools: traditional snapshot, Virtually Capacity-Free Snapshot (Vsnaps), and Virtually Instantaneous Snapclone. These data migration and replication tools have been refined in the EVA to provide tremendous flexibility and data protection.

Traditional snapshot

A traditional snapshot using EVA virtualization is similar to the traditional snapshot. In a standard snapshot, a set amount of capacity equal to the original volume is reserved for the snapshot. Data is not written into that reserved space until necessary. As the data changes in the original virtual disk, the data in the snapshot volume is updated with the original data. Even though the snapshot volume might not utilize all the reserved capacity, it must be set aside in reserve.

Virtually Capacity-Free Snapshot

In a Virtually Capacity-Free Snapshot, the storage system does not reserve capacity for the snapshot volume in advance. Rather, space in the snapshot volume is used only as the data of the original virtual disk is changed. This is called demand allocated. The snapshot volume is a new virtual disk that initially shares the original virtual disk map entries. As the original virtual disk is written, free space is consumed as necessary to preserve the original contents of the snapshot.

Both traditional snapshot and Vsnaps use unshared segments as data is written. The difference is that standard snapshots pre-allocate all space needed to do those unsharing operations, while Vsnaps allocate space only as needed at the time the unsharing is to be performed. This form of snapshot is called virtually capacity-free because the original and snapshot virtual disks share storage for similar segments of blocks. Other than this difference, the two types of snapshots are the same. It is important to note that a Virtually Capacity-Free Snapshot can be created from either Vraid 1, or Vraid 5. In addition, the EVA offers Cross-Vraid Snapshots, in which the Vraid level can be changed when snapping within the same disk group. This capability allows for more efficient disk utilization because redundant data can be changed from Vraid 1 to Vraid 5.

Virtually Instantaneous Snapclone

Virtually Instantaneous Snapclone is an improved type of data cloning similar to a traditional clone because duplicate space is reserved. A complete copy of the original virtual disk is made as quickly as data transfer rates permit, resulting in two identical independent copies of the data in the shortest time possible.

There is an important difference between Virtually Instantaneous Snapclone and the traditional clone. With traditional array controllers, the clone copy is not available until the copy is complete. With Virtually Instantaneous Snapclone, the snapclone data can be accessed almost immediately. As the Virtually Instantaneous Snapclone is being created, the controller can access the original virtual disk for the data and keep track of what data has changed since the instant the Virtually Instantaneous Snapclone was taken. The benefit of a Virtually Instantaneous Snapclone is that users get a point-in-time clone of a virtual disk essentially immediately.

Cross-Vraid Snapclones are another option for improving disk utilization. The EVA can change the Vraid level when creating snapclones across virtual disk groups, which would provide the customer significant cost savings when creating redundant copies of data. Not only could customers change the Vraid level of the snapclone, but they could also snapclone from standard high-performance Fibre Channel drives to less costly, higher density FATA drives.

In general, Virtually Instantaneous Snapclone is the optimum tool for long-term preservation of a copy or a series of copies of a disk, while capacity-free Vsnaps or traditional snapshots are better for creating short-lived, frozen images for applications such as backup.

Highlights of virtualization at the controller level with the EVA include:

- High performance through load leveling
- Efficient, automated storage management
- Better utilization of capacity and assets
- A dynamic storage environment
- More effective utilization of storage space with Virtually Capacity-Free Snapshots
- Immediate access with Virtually Instantaneous Snapclone
- Highly effective storage deployment environment

EVA remote replication with HP StorageWorks Continuous Access EVA

HP StorageWorks Continuous Access EVA provides a storage-based SAN application that brings the power of EVA virtualization to the arena of remote replication. Continuous Access EVA is enhanced to perform remote replication and deliver high data availability and performance on Fibre Channel based campus, metropolitan, or continental SANs.

Continuous Access EVA features the highest level of storage data-protection capabilities to meet business continuity implementation goals with:

- Fast application recovery with minimal or no transaction loss
- Application recovery to a metropolitan, regional, or continental site
- Disaster-tolerant solutions to ensure business continuance and company survival
- Essential component of every business continuity planning and implementation
- Creation of disaster-tolerant copies of critical business data
- No single-point-of-failure solution to increase the availability of your customers' data
- Easy-to-use graphical user interface for managing data replication between arrays

Virtualization and tiered storage

The EVA enables the user to identify different types of data that require different levels of accessibility and throughput. Virtualization enables easily managed, tiered storage solutions that comply with long-term data retention requirements and can be accessed instantaneously.

The EVA provides additional efficiencies for the customer's existing storage infrastructure and provides for the implementation of different classes of storage based on the data availability requirement of the customer's environment. Through the use of the FATA disk drive, the EVA provides users a low-cost storage solution for the intermediate storage case in which the customer's data might need to be migrated quickly between low-cost and high-performance devices. The EVA brings the same virtualization features to the management of the FATA disks at no additional cost.

The ability to use disk drives in this way helps an organization create a more adaptive enterprise. Moreover, the lower cost of the technology in the same footprint and the ability to expand the use of an existing device enhances the investment in the EVA and its value-added software and increases operational efficiency.

Additional information on the FATA disk drive can be found in the white paper, "HP StorageWorks Enterprise Virtual Array (EVA) Storage Solutions: High-Performance and Low-Cost Disk Drives," at <http://h71028.www7.hp.com/ERC/downloads/5982-7353EN.pdf>.

The business value of virtualization

Companies continue to face challenges associated with storing data in today's business environment.

Today's business challenges include:

- Limited to no economic visibility
- Shrinking or flat IT budget
- Brutal competition in many industries
- Internal headcount constraints
- Holding to major strategic projects
- Looking for creative new approaches

The common problems that IT professionals face every day include:

- The costs associated with the purchase of additional storage, physical space and environmental costs
- The operational time to bring storage into service
- The administration cost for storage allocation
- The people with the skill to manage storage
- More demand for data integration all within a multivendor environment

All of these problems can be summarized in the root issues of time, cost (resources), and space (capacity).

With virtualization, these common problems are reduced. Time, cost, and space are all better utilized to provide real value to the business.

Time

With the architecture of the EVA, virtualization at the controller level offers great performance advantages. Features such as automatic load balancing when new capacity is added to the storage pool can yield significant performance improvements. This means administrators no longer need to micromanage data placement, data replication, and performance tuning.

This technology offers truly scalable capacity, configuration flexibility, and more consistent user service levels. All of these benefits enable IT administrators to decrease the time required to maintain and manage their storage, while increasing the power and availability of their storage environment.

Cost

Virtualization can reduce total cost of ownership and lower the effective price per megabyte through management cost savings and better utilization of assets. The EVA, with its virtualizing architecture, utilizes a common management point for each unit on the SAN. As a result of this centralized management, storage administrators can manage more storage.

Storage pools can be managed easily with features such as dynamic expansion, simplified storage allocation, and load leveling to better utilize resources. These features enable the same number of administrators to manage significantly more storage, which is a real cost savings in today's climate of limited resources.

Cost savings are also critical as storage requirements continue to grow. A business must only purchase the capacity that is needed today thanks to virtualization features such as dynamic expansion and Virtually Capacity-Free Snapshot. It is no longer necessary to buy two or three times the amount of storage actually needed to make provisions for future growth. The key benefits of the EVA allow a business to buy storage when it is needed and not tie up precious capital.

Space

Through the benefits of virtualization, the EVA enables users to increase the value of their storage. They can take advantage of new flexibility and capacity savings through tools such as Virtually Capacity-Free Snapshot, Virtually Instantaneous Snapclone, and Cross-Vraid Snapshots and Snapclones. The versatility and cost effectiveness of the EVA is reemphasized through the introduction of FATA disk drives without restrictions into the current hardware and software configuration, which enables customers to optimize resources and assets.

Virtually Capacity-Free Snapshot, for example, enables administrators to replicate a point-in-time copy of data with only minimal use of capacity, which means that a storage manager has essentially twice the capacity in which to replicate data for backups, data mining, and other capacity-intensive activities. Because the controllers in the EVA use minimal capacity for data replication—while allowing any level of redundancy to be replicated—the result is a much lower “effective” cost per megabyte because the user can access more storage for the money.

Due to the Fibre Channel virtualization architecture of the EVA, the amount of capacity behind the controller is significant. The EVA can support up to 240 drives behind one pair of controllers with the EVA8100 and 168 drives in one cabinet. The result is denser storage configurations in the same footprint of a standard 42U cabinet.

The elimination of stranded and over-allocated capacity through virtualization is another important benefit. Traditionally, as data is moved and changed over time, segments of unused or “stranded” capacity—that is, storage which is allocated to an application or server but does not get used—can result. This condition leads to wasted capacity and resources. In addition, because of the difficulty in expanding raidsets in traditional storage, the raidsets would typically be created with significantly more capacity than needed at the time to prepare for future expansion. However, the EVA can move this stranded or over-allocated capacity into an available storage pool. As a result of this functionality, the user can once again use more of their storage capacity.

Conclusion

While virtualization is being addressed at all levels of the SAN, the most powerful application of this technology today is at the storage system level. With the HP StorageWorks Enterprise Virtual Array, organizations can take advantage of the many benefits of virtualization with a solution designed specifically for open systems.

Virtualization at the storage system level with the EVA offers many advantages. The common problems of storage environments are all curtailed. Cost, space, and management resources are all better utilized in a virtualized environment, providing real value to business.

The EVA provides time savings through automatic performance tuning as well as a high controller throughput. As a result, administrators no longer need to micromanage data placement and performance tuning. The result is less storage downtime.

Through centralized management, virtualization helps reduce costs. The same number of storage administrators can manage more storage while company assets, such as disk space, are better utilized.

IT professionals benefit from the flexibility and capacity savings through support for virtual disk mirrors, Virtually Capacity-Free Snapshot, and Virtually Instantaneous Snapclone. These tools deliver optimized capacity and greater flexibility to storage administrators.

Figure 5. The value of EVA virtualization

Lower management and training costs <ul style="list-style-type: none">• Easy to use, intuitive interface• Easy to create Vraid volumes (LUNs)• Unifies storage into a common pool	Better application availability <ul style="list-style-type: none">• Dynamic LUN expansion• Enterprise-class availability• Dynamic expansion• No storage reconfiguration downtime
Buy less upfront and long-term <ul style="list-style-type: none">• Better capacity utilization• Almost eliminate stranded capacity• Virtual capacity free snapshots save disk space	Improved performance and serve more customers <ul style="list-style-type: none">• Striping across all drives in the disk group reduces hotspots and bottlenecks• Automatic load leveling

As an ideal storage system for SAN environments, the HP StorageWorks Enterprise Virtual Array offers true value in today's demanding storage environments. IT managers who want to save money, space, and time will find the EVA the right solution to meet their business-critical storage needs.

For more information

- Learn more about HP StorageWorks virtualization at www.hp.com/go/StorageVirtualization
- You can learn more about the HP StorageWorks EVA family of products at www.hp.com/go/EVA
- Read the report by the Edison Group that determined a storage administrator can perform a series of standard tasks in 70 percent fewer steps on an HP StorageWorks EVA than on “traditional” disk array systems. Find it at <http://h71028.www7.hp.com/ERC/downloads/4AA1-6634ENW.pdf>
- You can find more white papers about the EVA and other HP StorageWorks array solutions at <http://h18006.www1.hp.com/storage/arraywhitepapers.html>
- Try the HP StorageWorks product selector tool to see which product is right for you at <http://www.hp.com/go/product-selector>

© Copyright 2008 Hewlett-Packard Development Company, L.P. The information contained herein is subject to change without notice. The only warranties for HP products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty. HP shall not be liable for technical or editorial errors or omissions contained herein.

Linux is a U.S. registered trademark of Linus Torvalds. Microsoft and Windows are U.S. registered trademarks of Microsoft Corporation. UNIX is a registered trademark of The Open Group.

4AA1-8130ENW, February 2008

